

Computational Geometry - HW1

Shaza Abdallah - 323032474

November 22, 2025

1 Question 1

1.1 (a)

Proof. Let S_1 and S_2 be convex sets, and assume that $S_1 \cap S_2 \neq \emptyset$. Take any two points $p, q \in S_1 \cap S_2$. Then $p, q \in S_1$ and $p, q \in S_2$. Since S_1 is convex, the entire segment pq lies in S_1 ; similarly, since S_2 is convex, the segment pq lies in S_2 . Thus every point of pq lies in both sets, and so $pq \subseteq S_1 \cap S_2$. Therefore, $S_1 \cap S_2$ is convex.

1.2 (b)

Counterexample.

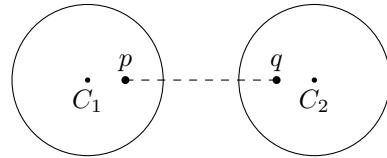


Figure 1: Two disjoint circles C_1 and C_2 .

Let C_1 and C_2 be two circles in \mathbb{R}^2 whose interiors and boundaries do not intersect, and consider the set

$$S = C_1 \cup C_2.$$

We show that S is not convex.

Pick a point $p \in C_1$ and a point $q \in C_2$, for example as in the figure above. By definition of convexity, if S were convex, then the entire line segment pq would have to be contained in S .

However, since the circles are disjoint and separated in space, the segment pq contains points that lie strictly between the two circles, hence outside both C_1 and C_2 . Therefore these points are not in $S = C_1 \cup C_2$.

1.3 (c)

Counterexample.

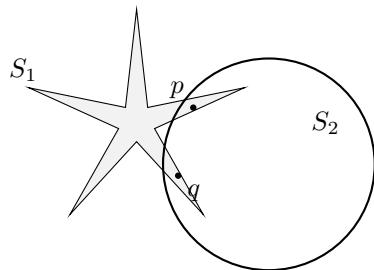


Figure 2: A star-shaped set S_1 and a circle S_2 .

The intersection

$$S = S_1 \cap S_2$$

consists of two disjoint connected components, one at the top of the circle and one at the bottom. Assume for contradiction that S is star-shaped with respect to some point $x \in S$. Then p lies in exactly one of the two components. Choose a point q in the other component. Any line segment from p to q must leave S , so in particular $pq \not\subseteq S$, contradicting the definition of a star-shaped set. Therefore $S_1 \cap S_2$ is not star-shaped.

1.4 (d)

Counterexample. Consider the same sets S_1 and S_2 as in part (c): S_1 is the star-shaped set (the star), and S_2 is the circle, which is convex. Their intersection

$$S = S_1 \cap S_2$$

consists of two disjoint connected components.

Pick points p and q in different components of S (as indicated in the figure). The line segment $[p, q]$ necessarily passes through points that lie outside S , so $pq \not\subseteq S$. Hence S is not convex.

2 Question 2

Let S be a set of n circles in the plane, each given by the endpoints of its horizontal diameter. We describe a plane-sweep algorithm that computes all k intersection points in $O((n + k) \log n)$ time.

Events

We use a priority queue sorted by x -coordinate. There are three event types:

1. *Circle Start*: the leftmost point $(x_i - r_i, y_i)$ of circle C_i .
2. *Circle End*: the rightmost point $(x_i + r_i, y_i)$ of C_i .
3. *Circle Intersect*: an intersection point of two circles.

Sweep-line Status

Unlike line segments, a circle does not have a single well-defined vertical order along the sweep line. At a sweep position $x = x_0$, circle C_i intersects the vertical line in two points:

$$y_i \pm \sqrt{r_i^2 - (x_0 - x_i)^2}.$$

Thus the sweep-line status stores *two* entries for each active circle:

$$\text{TopHalf}(C_i), \quad \text{BottomHalf}(C_i),$$

each ordered by its y -coordinate on the sweep line. The status structure is a balanced BST ordered by these y -values.

Handling Events

1. Circle Start. When encountering the leftmost point of C_i :

1. Insert $\text{TopHalf}(C_i)$ and $\text{BottomHalf}(C_i)$ into the status.
2. For each of them, check for intersections with the immediate predecessor and successor arcs in the BST.
3. For every real intersection point whose x -coordinate is in the future, insert a *Circle Intersect* event.

2. Circle End. When reaching the rightmost point of C_i :

1. Let $a = \text{BottomHalf}(C_i)$ and $b = \text{TopHalf}(C_i)$.
2. Before removing a and b , let p be the predecessor of a and let q be the successor of b .
3. Remove a and b from the status.
4. Check whether p and q intersect; if so, add the corresponding intersection event.

3. Circle Intersect. Suppose arc A and arc B intersect at (x^*, y^*) .

1. Output the intersection point.
2. In the status, swap the order of A and B (only the two semicircles swap).
3. After the swap:
 - Check for intersections between A and its new predecessor and successor.
 - Check for intersections between B and its new predecessor and successor.
4. Insert any newly discovered valid intersection events.

Running Time

There are $2n$ start/end events and k intersection events. Each event triggers $O(1)$ neighbor checks and $O(\log n)$ updates in the balanced BST and event queue. Therefore the total running time is

$$O((n + k) \log n),$$

3 Question 3

3.1 (a)

- $\text{Twin}(\text{Twin}(e)) = e$ (always true)

Each undirected edge in the DCEL is represented by exactly two half-edges, which are twins of one another and point in opposite directions.

- $\text{Next}(\text{Prev}(e)) = e$ (always true)

For any face, its boundary is stored as a circular doubly linked list of half-edges. By definition of a doubly linked list,

$$\text{Next}(\text{Prev}(e)) = e$$

- $\text{Twin}(\text{Prev}(\text{Twin}(e))) = \text{Next}(e)$ (not always true)

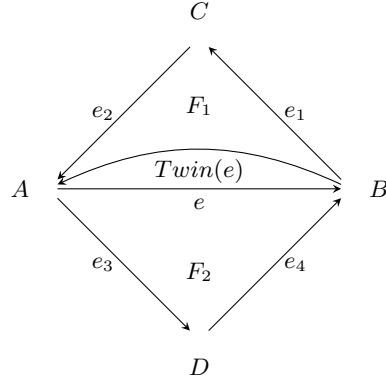


Figure 3: Two triangles sharing edge \overline{AB} , showing $e : A \rightarrow B$ and its twin $\text{Twin}(e) : B \rightarrow A$ as distinct half-edges. On face F_2 , the boundary order is $D \rightarrow B \rightarrow A \rightarrow D$, so $\text{Prev}(\text{Twin}(e)) = e_4 : D \rightarrow B$, and therefore $\text{Twin}(\text{Prev}(\text{Twin}(e))) = B \rightarrow D \neq \text{Next}(e) = B \rightarrow C$.

- $\text{IncidentFace}(e) = \text{IncidentFace}(\text{Next}(e))$ (always true)
All half-edges around the boundary of the same face share the same IncidentFace pointer. Since $\text{Next}(e)$ walks along the same face,

$$\text{IncidentFace}(e) = \text{IncidentFace}(\text{Next}(e)).$$

3.2 (b)

- List all vertices that are connected by an edge to a given vertex v

Algorithm 1 ListNeighbors(v)

```

1:  $e \leftarrow \text{IncidentEdge}(v)$ 
2:  $d \leftarrow e$ 
3: repeat
4:    $S.\text{add}(\text{origin}(\text{Twin}(d)))$ 
5:    $d \leftarrow \text{Next}(\text{Twin}(d))$ 
6: until  $d = e$ 
7: return  $S$ 

```

- List all edges that bound a given face f in a not necessarily connected subdivision

Algorithm 2 ListEdges(f)

```

1:  $S \leftarrow \emptyset$ 
2:  $e_{\text{out}} \leftarrow \text{OuterComponent}(f)$ 
3: if  $e_{\text{out}} \neq \text{null}$  then
4:   CountCycle( $e_{\text{out}}, S$ )
5: end if
6: for each  $e_{\text{in}}$  in  $\text{InnerComponents}(f)$  do
7:   CountCycle( $e_{\text{in}}, S$ )
8: end for
9: return  $S$ 

```

Algorithm 3 CountCycle(e, S)

```

1:  $d \leftarrow e$ 
2: repeat
3:    $S.\text{add}(d)$ 
4:    $d \leftarrow \text{Next}(d)$ 
5: until  $d = e$ 

```

- List all faces that have at least one vertex on the outer boundary of the subdivision

Algorithm 4 ListFacesTouchingOuterBoundary()

```
1:  $S \leftarrow \emptyset$ 
2: for each face  $f$  in the subdivision do
3:   if  $OuterComponent(f) = \text{nil}$  then
4:      $OuterFace \leftarrow f$ 
5:     break
6:   end if
7: end for
8:  $BoundaryVertices \leftarrow \emptyset$ 
9: for each  $e_0$  in  $InnerComponents(OuterFace)$  do
10:    $d \leftarrow e_0$ 
11:   repeat
12:      $BoundaryVertices.add(origin(d))$ 
13:      $d \leftarrow Next(d)$ 
14:   until  $d = e_0$ 
15: end for
16: for each vertex  $v$  in  $BoundaryVertices$  do
17:    $h \leftarrow IncidentEdge(v)$ 
18:    $g \leftarrow h$ 
19:   repeat
20:      $f \leftarrow IncidentFace(g)$ 
21:      $S.add(f)$ 
22:      $g \leftarrow Next(Twin(g))$ 
23:   until  $g = h$ 
24: end for
25: return  $S$ 
```

3.3 (d)

The subdivision can have one face. In a DCEL we always have

$$IncidentFace(e) = IncidentFace(Next(e)).$$

Under the given assumption $Twin(e) = Next(e)$ for every half-edge e , we obtain

$$IncidentFace(e) = IncidentFace(Next(e)) = IncidentFace(Twin(e)).$$

Thus a half-edge and its twin always belong to the *same* face for all edges.

4 Question 4

4.1 (a)

Let $P = (p_0, p_1, \dots, p_{n-1})$ be a simple polygon listed in cyclic order, and let $\ell : y = mx + b$ be the line with respect to which we want to test monotonicity.

Define the projection function

$$h(x, y) = x + my.$$

This value gives the position of a point along the direction of ℓ ; points lying on any line perpendicular to ℓ have equal h -value.

Algorithm:

1. for each vertex $p_i = (x_i, y_i)$ of P , compute

$$h_i = h(p_i) = x_i + my_i.$$

2. find the vertex p_s having the minimum projection value and the vertex p_t having the maximum projection value:

$$s = \arg \min_i h_i, \quad t = \arg \max_i h_i.$$

These serve as the “lowest” and “highest” vertices of the polygon in the direction of ℓ .

3. starting at p_s , follow the polygon boundary forward until reaching p_t . Check that the sequence

$$h_s, h_{s+1}, h_{s+2}, \dots, h_t$$

is monotonically nondecreasing. If a decrease occurs (i.e. $h_{i+1} < h_i$ for some consecutive pair), then P is not ℓ -monotone.

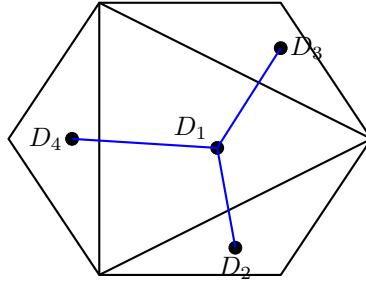
4. again start from p_s , but now follow the vertices backward (decreasing the index modulo n) until reaching p_t . Confirm that the projection sequence along this second chain is also monotonically nondecreasing. A decrease anywhere implies the polygon is not ℓ -monotone.
5. if both boundary chains from p_s to p_t have nondecreasing projection values, then the polygon is ℓ -monotone. Otherwise, it is not.

The function $h(x, y) = x + my$ increases exactly in the direction parallel to ℓ . Lines perpendicular to ℓ are precisely those on which h is constant. Thus, checking monotonicity with respect to ℓ amounts to verifying that along each boundary chain between the minimal and maximal projection vertices, the polygon never moves “backwards” in h . A decrease in projection means that some line perpendicular to ℓ intersects the polygon in two disjoint segments. Conversely, if both chains have nondecreasing projections, every perpendicular line intersects P in at most one connected interval, which is exactly the definition of ℓ -monotonicity.

Computing all projection values takes $O(n)$ time, finding the extremal vertices takes $O(n)$ time, and traversing both boundary chains also takes $O(n)$. Thus, the full algorithm runs in $O(n)$

4.2 (b)

counterexample.



This polygon is convex, and therefore monotone with respect to every line. However, in the shown triangulation the central triangle corresponds to a dual vertex D_1 that is adjacent to three other triangles. Thus $\deg(D_1) = 3$, so the dual graph is not a chain.

5 Question 5

5.1 (a)

Proof. Let P have h holes and n vertices in total. We prove the theorem by induction on h primarily, and on n secondarily.

Base case. For $h = 0$, P is a simple polygon with no holes. It was proved in the lecture that every simple polygon can be triangulated. Thus the claim holds for $h = 0$.

Induction hypothesis. Assume the theorem holds for every polygon P_1 with $h_1 < h$ holes, and also holds for every polygon P_2 with $h_2 \leq h$ holes and strictly fewer $n_2 < n$ vertices.

Induction step. Let P be a polygon with h holes and n vertices. Choose an arbitrary convex vertex v_2 on the outer boundary of P , and let v_1 and v_3 be its two neighbours on the boundary.

Consider the segment v_1v_3 . If the open segment v_1v_3 lies entirely inside P , then $d = v_1v_3$ is an internal diagonal of P . Otherwise, let $d = v_2x$, where x is the closest vertex to v_2 measured along the line perpendicular to v_1v_3 .

We distinguish two cases.

Case 1. The diagonal d has one endpoint on a hole. Cutting along d decreases the number of holes by 1, but increases the number of vertices by 2. Thus the resulting polygon has $(h - 1)$ holes, and hence the induction hypothesis on h applies. Therefore it can be triangulated, and this triangulation is also a triangulation of P .

Case 2. Both endpoints of d lie on the outer boundary of P . Then d partitions P into two polygons P_1 and P_2 , each with at most h holes and strictly fewer than n vertices. By the induction hypothesis on n , both P_1 and P_2 admit triangulations, which together form a triangulation of P . \square

5.2 (b)

Proof. We prove the statement by induction on the pair (h, n) , with induction on h primarily and on n secondarily.

Base case $h = 0$. If P is a simple polygon with no holes, then any triangulation has exactly $n - 2$ triangles, as proved in the lecture. Thus the formula holds for $h = 0$:

$$n - 2 + 2h = n - 2.$$

Induction hypothesis. Assume that the statement holds for every polygon P_1 with $h_1 < h$ holes, and for every polygon P_2 with $h_2 = h$ holes and $n_2 < n$ vertices.

Induction step. Let P have h holes and n vertices. Choose a diagonal d obtained by the same construction as in the proof of triangulability: either $d = v_1v_3$ or $d = v_2x$, depending on whether v_1v_3 is internal. We again consider two cases.

Case 1. The diagonal d connects the outer boundary to a hole. Cutting along d produces a polygon P' with $h - 1$ holes and $n + 2$ vertices. By the induction hypothesis on h , every triangulation of P' contains

$$(n + 2) - 2 + 2(h - 1) = n - 2 + 2h$$

triangles. Since every triangulation of P corresponds to a triangulation of P' , the same number of triangles appears in P .

Case 2. Both endpoints of d lie on the outer boundary of P . Then d partitions P into two polygons P_1 and P_2 with $n_1 + n_2 = n + 2$ and $h_1 + h_2 = h$, and with $n_1, n_2 < n$. By the induction hypothesis on n , their triangulations contain

$$n_1 - 2 + 2h_1 \quad \text{and} \quad n_2 - 2 + 2h_2$$

triangles, respectively. Adding these counts and observing that

$$(n_1 - 2 + 2h_1) + (n_2 - 2 + 2h_2) = n - 2 + 2h$$

gives the desired total number of triangles in P . \square

5.3 (c)

Let T_n denote the number of triangulations of a convex polygon with n vertices. Fix the edge p_1p_n . In every triangulation, this edge must belong to a unique triangle $\triangle p_1p_ip_n$ for some i with $2 \leq i \leq n - 1$.

This choice splits the polygon into two smaller convex polygons:

$$\{p_1, p_2, \dots, p_i\} \quad \text{and} \quad \{p_i, p_{i+1}, \dots, p_n\},$$

having i and $n - i + 1$ vertices respectively. Their triangulations are independent, so the number of triangulations with triangle $p_1p_ip_n$ is

$$T_i \cdot T_{n-i+1}.$$

Summing over all possible i gives the recurrence

$$T_n = \sum_{i=2}^{n-1} T_i T_{n-i+1},$$

with the base case $T_2 = 1$ (a polygon with two vertices has one trivial “triangulation”).